# CSE 305: Language Interpreter Design

Part 1 Due: Friday November 2 2018, at 11:59 pm
Part 2 Due: Sunday November 18 2018, at 11:59 pm
Part 3 Due: Tuesday December 4 2018, at 11:59 pm

## 1   Overview

The goal of this homework is to understand and build an interpreter in two languages (you may choose between Python and Java, and you must use OCaml) for a small OCaml-like, stack based, bytecode language. The project is broken down into three parts. Part 1 is defined in Section 4, Part 2 is defined in Section 5, and Part 3 is defined in Section 6. Each part is worth 100 points, 50 for each language. You should spend roughly two weeks for each part. Test cases for each part will be provided in Autolab. Put your answers for Java/Python and OCaml in files named, respectively:

1. interpreter.java

2. interpreter.py

3. interpreter.ml

These files should contain a function (static method in Java) called `interpreter` that takes two strings (`interpreter(input, ouput)`). You can submit multiple files as a tar or zip archive through Autolab. You will submit one solution for each separate part. Each part is graded individually and a solution will be provided after the submission deadline. Late submissions will not be accepted and will be given a score of 0. Test cases will also be provided on Piazza for you to test your code locally. These will not be exhaustive, so you are highly encouraged to write your own tests to check your interpreter against all the functionality described in this document.

## 2   Functionality

Your interpreter function (or static method) should take in two arguments, the file you are reading from (input) and the file name of your output file (output).

```
val interpreter :  string -> string -> unit = <fun>
public static void interpreter (String inputFile, String outputFile){}
def interpreter(inputFile, outputFile):
```

Input and output will be passed in as strings that represent paths to files just like in your first homework assignment. Your function should write to the output file the contents of the final stack your interpreter produces. In the examples below, the input file is read from top to bottom and each command is executed by your interpreter in the order it was read. You may find it useful to read in all of the commands into a list or other data structure prior to executing them. The input file can be arbitrarily long.

| input | stack |
|---|---|
| push 1 | 1 |
| quit | |

| input | stack |
|---|---|
| push 6 | |
| push 2 | |
| div | :error: |
| mul | 3 |
| quit | |

| input | stack |
|---|---|
| push 5 | |
| neg | |
| push 10 | 30 |
| push 20 | -5 |
| add | |
| quit | |

| input | stack |
|---|---|
| push :true: | |
| push 7 | |
| push 8 | |
| push :false: | -1 |
| pop | :true: |
| sub | |
| quit | |

| input | stack |
|---|---|
| push 10 | |
| push 2 | |
| push 8 | |
| mul | 23 |
| add | |
| push 3 | |
| sub | |
| quit | |

# 3 Grammar

The following is a context free grammar for the bytecode language you will be implementing. Terminal symbols are identified by `monospace font`, and nonterminal symbols are identified by *italic font.* Anything enclosed in ⟨wide-angle brackets⟩ denotes an optional character (zero or one occurrences). The form '( $set_1$ | $set_2$ | $set_n$ )' means a choice of one character from any one of the $n$ sets. A set with a superscript asterisk means zero or more occurrences of characters in that set, and a superscript plus-sign ($+$) means one or more occurrences.

The set *digit* is the set of digits $\{0,1,2,3,4,5,6,7,8,9\}$, *letter* is the set of all characters in the English alphabet (lowercase and uppercase), and *ASCII* is the ASCII character set. The set *simpleASCII* is *ASCII* without quotation marks and the backslash character.

## 3.1 Constants

*const* ::= *int* | *bool* | *error* | *string* | *name*

*int* ::= ⟨−⟩ *digit*$^+$

*bool* ::= `:true:` | `:false:`

*error* ::= `:error:`

*string* ::= `"`*simpleASCII*$^+$`"`

*simpleASCII* ::= *ASCII* \ {`'\'`, `'"'`}

*name* ::= ⟨_⟩ *letter* (*letter* | *digit* | _)$^*$

## 3.2 Programs

*prog* ::= *com*$^+$ `quit`

*com* ::= `push` *const* | `add` | `sub` | `mul` | `div` | `rem` | `neg` | `swap` | `pop` | `cat` | `and` | `or` | `not`
    | `lessThan` | `equal` | `if` | `bind` | `let` *com*$^+$ `end` | *funBind* *com*$^+$ ⟨`return`⟩ `funEnd` | `call`

*funBind* ::= (`fun` | `inOutFun`) *name*$_1$ *name*$_2$

# 4 Part 1: Basic Computation    Due: Friday November 2 2018, at 11:59 pm

Your interpreter should be able to handle the following commands:

## 4.1 push

### 4.1.1 Pushing Integers to the Stack

$$\text{push } num$$

where *num* is an integer possibly with a '-' suggesting a negative value. Here, '-0' should be regarded as '0'. Entering this expression will simply push *num* onto the stack. For example,

| input | stack |
|---|---|
| push 5 | 0 |
| push -0 | 5 |

If *num* is not an integer, only push the error literal (:error:) onto the stack instead of pushing *num*. For example,

| input | stack |
|---|---|
| push 5 | :error: |
| push 2.5 | :error: |
| push -x | 5 |

### 4.1.2 Pushing Strings to the Stack

$$\text{push } string$$

where *string* is a string literal consisting of a sequence of characters enclosed in double quotation marks, as in "this is a string". Executing this command would push the string onto the stack:

| input | stack |
|---|---|
| push "deadpool" | this this a string |
| push "batman" | batman |
| push "this is a string " | deadpool |

Spaces are preserved in the string, ie any preceding or trailing whitespace must be kept inside the string that is pushed to the stack:

| input | stack |
|---|---|
| push " deadp ool " | this␣this␣a␣string␣␣ |
| push "this is a string " | ␣deadp␣ool␣ |

You can assume that the string value would always be legal and not contain quotations or escape sequences within the string itself, i.e double quotes nor backslashes will appear inside a string.

## 4.2  Pushing Names to the Stack

$$\text{push } name$$

where *name* consists of a sequence of letters, digits, and underscores, starting with a letter or underscore.

1. example

| input | | stack | | stack |
|---|---|---|---|---|
| push a | $\rightarrow$ | a | $\rightarrow$ | 13 |
| push 13 | | | | a |

2. example

| input | | stack | | stack |
|---|---|---|---|---|
| push __name1__ | $\rightarrow$ | __name1__ | $\rightarrow$ | 3 |
| push 3 | | | | __name1__ |

To bind 'a' to the value 13 and name1 to the value 3, we will use 'bind' operation which we will see later (Section 5.7) You can assume that name will not contain any illegal tokens—no commas, quotation marks etc. It will always be a sequence of letters and digits starting with a letter.

## 4.3  pop

Remove the top value from the stack. If the stack is empty, an error literal (`:error:`) will be pushed onto the stack. For example,

| input | | stack | | stack | | input |
|---|---|---|---|---|---|---|
| push 5 | $\rightarrow$ | 5 | $\rightarrow$ | | $\rightarrow$ | :error: |
| pop | | | | | | |
| pop | | | | | | |

## 4.4  boolean

There are two kinds of boolean literals: `:true:` and `:false:`. Your interpreter should push the corresponding value onto the stack. For example,

| input | stack |
|---|---|
| push 5 | :true: |
| push :true: | 5 |

## 4.5  error

Similar with boolean literals, pushing error literal will push `:error:` onto the stack.

## 4.6 add

add refers to integer addition. Since this is a binary operator, it consumes the top two values in the stack, calculate sum and push the result back to the stack. If one of the following cases occurs, which means there is an error, any values popped out from the stack should be pushed back in the same order, then a value :error: should also be pushed onto the stack:

- not all top two values are integer numbers

- only one value in the stack

- stack is empty

for example,

| input |
|-------|
| push 5 |
| push 8 |
| add |

$\rightarrow$

| stack |
|-------|
| 8 |
| 5 |

$\rightarrow$

| stack |
|-------|
| 13 |

For another example, if there is only one number in the stack and we use add, an error will occur. Then 5 should be pushed back as well as :error:

| input |
|-------|
| push 5 |
| add |

$\rightarrow$

| stack |
|-------|
| 5 |

$\rightarrow$

| stack |
|-------|
| :error: |
| 5 |

## 4.7 sub

The command sub refers to integer subtraction. It is a binary operator and works in the following way:

- if top two elements in the stack are integer numbers, pop the top element(y) and the next element(x), subtract y from x, and push the result x-y back onto the stack

- if the top two elements in the stack are not all integer numbers, push them back in the same order and push :error: onto the stack

- if there is only one element in the stack, push it back and push :error: onto the stack

- if the stack is empty, push :error: onto the stack

for example,

| input |
|-------|
| push 5 |
| push 8 |
| sub |

$\rightarrow$

| stack |
|-------|
| 8 |
| 5 |

$\rightarrow$

| stack |
|-------|
| -3 |

For another example, if one of the top two values in the stack is not a numeric number when sub is used, an error will occur. Then 5 and :false: should be pushed back as well as :error:

| input |
|-------|
| push 5 |
| push :false: |
| sub |

$\rightarrow$

| stack |
|-------|
| 5 |

$\rightarrow$

| stack |
|-------|
| :false: |
| 5 |

$\rightarrow$

| stack |
|-------|
| :error: |
| :false: |
| 5 |

## 4.8  mul

The command mul refers to integer multiplication. It is a binary operator and works in the following way:

- if top two elements in the stack are integer numbers, pop the top element(y) and the next element(x), multiply x by y, and push the result x*y back onto the stack

- if the top two elements in the stack are not all integer numbers, push them back in the same order and push :error: onto the stack

- if there is only one element in the stack, push it back and push :error: onto the stack

- if the stack is empty, push :error: onto the stack

For example:

| input |
|-------|
| push 5 |
| push 8 |
| mul |

→

| stack |
|-------|
| 8 |
| 5 |

→

| stack |
|-------|
| 40 |

If the stack empty when mul is executed, an error will occur and :error: should be pushed onto the stack:

| input |
|-------|
| mul |

→

| stack |
|-------|

→

| stack |
|-------|
| :error: |

## 4.9  div

The command div refers to integer division. It is a binary operator and works in the following way:

- if top two elements in the stack are integer numbers, pop the top element(y) and the next element(x), divide x by y, and push the result $\frac{x}{y}$ back onto the stack

- if top two elements in the stack are integer numbers but y equals to 0, push them back in the same order and push :error: onto the stack

- if the top two elements in the stack are not all integer numbers, push them back in the same order and push :error: onto the stack

- if there is only one element in the stack, push it back and push :error: onto the stack

- if the stack is empty, push :error: onto the stack

For example:

| input |
|-------|
| push 5 |
| push 8 |
| div |

→

| stack |
|-------|
| 8 |
| 5 |

→

| stack |
|-------|
| 0 |

If the top element in the stack equals to 0, there will be an error if div is executed. In such situations 5 and 0 should be pushed back onto the stack as well as :error:

| input |
|-------|
| push 5 |
| push 0 |
| div |

→

| stack |
|-------|
| 0 |
| 5 |

→

| stack |
|-------|
| :error: |
| 0 |
| 5 |

## 4.10   rem

The command rem refers to the remainder of integer division. It is a binary operator and works in the following way:

- if top two elements in the stack are integer numbers, pop the top element(y) and the next element(x), calculate the remainder of $\frac{x}{y}$, and push the result back onto the stack

- if top two elements in the stack are integer numbers but y equals to 0, push them back in the same order and push :error: onto the stack

- if the top two elements in the stack are not all integer numbers, push them back and push :error: onto the stack

- if there is only one element in the stack, push it back and push :error: onto the stack

- if the stack is empty, push :error: onto the stack

For example,

| input |
|-------|
| push 5 |
| push 8 |
| rem |

$\rightarrow$

| stack |
|-------|
| 8 |
| 5 |

$\rightarrow$

| stack |
|-------|
| 5 |

If one of the top two elements in the stack is not an integer, an error will occur if rem is executed. If this occurs the top to elements should be pushed back onto the stack as well as :error:. For example:

| input |
|-------|
| push 5 |
| push :false: |
| rem |

$\rightarrow$

| stack |
|-------|
| :false: |
| 5 |

$\rightarrow$

| stack |
|-------|
| :error: |
| :false: |
| 5 |

## 4.11   neg

The command neg is to calculate the negation of an integer (negation of 0 should still be 0). It is unary therefore consumes only the top element from the stack, calculate its negation and push the result back. A value :error: will be pushed onto the stack if:

- the top element is not an integer, push the top element back and push :error:

- the stack is empty, push :error: onto the stack

For example:

| input |
|-------|
| push 5 |
| neg |

$\rightarrow$

| stack |
|-------|
| 5 |

$\rightarrow$

| stack |
|-------|
| -5 |

If the value on top of the stack is not an integer, when neg is used, that value should be pushed back onto the stack as well as :error:. For example:

| input |
|-------|
| push 5 |
| neg |
| push :true: |
| neg |

$\rightarrow$

| stack |
|-------|
| -5 |

$\rightarrow$

| stack |
|-------|
| :true: |
| -5 |

$\rightarrow$

| stack |
|-------|
| :error: |
| :true: |
| -5 |

## 4.12   swap

The command swap interchanges the top two elements in the stack, meaning that the first element becomes the second and the second becomes the first. A value :error: will be pushed onto the stack if:

- there is only one element in the stack, push the element back and push :error:

- the stack is empty, push :error: onto the stack

For example:

| input |
|-------|
| push 5 |
| push 8 |
| push :false: |
| swap |

→

| stack |
|-------|
| 8 |
| 5 |

→

| stack |
|-------|
| :false: |
| 8 |
| 5 |

→

| stack |
|-------|
| 8 |
| :false: |
| 5 |

If there is only one element in the stack when swap is used, an error will occur and :error: should be pushed onto the stack. Now we have two elements in the stack (5 and :error:), therefore the second swap will interchange the two elements:

| input |
|-------|
| push 5 |
| swap |
| swap |

→

| stack |
|-------|
| 5 |

→

| stack |
|-------|
| :error: |
| 5 |

→

| stack |
|-------|
| 5 |
| :error: |

## 4.13   quit

The command quit causes the interpreter to stop. Then the whole stack should be printed out to an output file that is specified as the second argument to the interpret function.

# 5   Part 2: Variables and Scope     Due Date: Sunday November 18 2018, at 11:59 pm

In part 2 of the interpreter you will be expanding the types of computation you will be able to perform, adding support for immutable variables, and structures for expressing scope.

## 5.1   cat

The cat command computes the concatenation of the top two elements in the stack and pushes the result onto the stack. The top two values of the stack—x and y—are popped off and the result is the string y concatenated with x.
:error: will be pushed onto the stack if:

- there is only one element in the stack, push the element back and push :error:

- the stack is empty, push :error: onto the stack

- if either of the top two elements are not strings, push the elements back onto the stack, and then push :error:

– NB: Recall that names and strings are different.

Example:

| input |
|---|
| push "hello " |
| push "world!" |
| cat |

→

| stack |
|---|
| hello |

→

| stack |
|---|
| world! |
| hello |

→

| stack |
|---|
| hello world! |

Note that strings can contain spaces, punctuation marks, and other special characters. You may assume that strings only contain ASCII characters and have no escape sequences, eg \n and \t.

| input |
|---|
| push "Michael" |
| push Scott |
| cat |

→

| stack |
|---|
| Michael |

→

| stack |
|---|
| Scott |
| Michael |

→

| stack |
|---|
| :error: |
| Scott |
| Michael |

## 5.2   and

The command and performs the logical conjunction of the top two elements in the stack and pushes the result (a single value) onto the stack.
:error: will be pushed onto the stack if:

- there is only one element in the stack, push the element back and push :error:

- the stack is empty, push :error: onto the stack

- if either of the top two elements arenâĂŹt Boolean, push back the elements and push :error:

For example:

| input |
|---|
| push :true: |
| push :false: |
| and |

→

| stack |
|---|
| :true: |

→

| stack |
|---|
| :false: |
| :true: |

→

| stack |
|---|
| :false: |

Consider another example:

| input |
|---|
| push :true: |
| and |

→

| stack |
|---|
| :true: |

→

| stack |
|---|
| :error: |
| :true: |

## 5.3   or

The command or performs the logical disjunction of the top two elements in the stack and pushes the result (a single value) onto the stack.
:error: will be pushed onto the stack if:

- there is only one element in the stack, push the element back and push :error:

- the stack is empty, push :error: onto the stack

- if either of the top two elements aren't Boolean, push back the elements and push :error:

For example:

```
┌─────────────┐       ┌────────┐       ┌────────┐       ┌────────┐
│    input    │       │ stack  │       │ stack  │       │ stack  │
│ push :true: │   →   │ :true: │   →   │ :false:│   →   │ :true: │
│ push :false:│       └────────┘       │ :true: │       └────────┘
│ or          │                        └────────┘
└─────────────┘
```

Consider another example:

```
┌──────────────────┐       ┌────────┐       ┌──────────┐       ┌──────────┐
│      input       │       │ stack  │       │  stack   │       │  stack   │
│ push :false:     │   →   │ :false:│   →   │ khaleesi │   →   │ :error:  │
│ push "khaleesi"  │       └────────┘       │ :false:  │       │ khaleesi │
│ or               │                        └──────────┘       │ :false:  │
└──────────────────┘                                           └──────────┘
```

## 5.4  not

The command not performs the logical negation of the top element in the stack and pushes the result (a single value)onto the stack. Since the operator is unary, it only consumes the top value from the stack. The :error: value will be pushed onto the stack if:

- the stack is empty, push :error: onto the stack

- if the top element isn't Boolean, push back the element and push :error:

For example:

```
┌─────────────┐       ┌────────┐       ┌────────┐
│    input    │       │ stack  │       │ stack  │
│ push :true: │   →   │ :true: │   →   │ :false:│
│ not         │       └────────┘       └────────┘
└─────────────┘
```

Consider another example:

```
┌─────────┐       ┌────────┐       ┌─────────┐
│  input  │       │ stack  │       │  stack  │
│ push 3  │   →   │   3    │   →   │ :error: │
│ not     │       └────────┘       │   3     │
└─────────┘                        └─────────┘
```

## 5.5  equal

The command equal refers to numeric equality (so you are not supporting string comparisons). This operator consumes the top two values on the stack and pushes the result (a single boolean value) onto the stack. The :error: value will be pushed onto the stack if:

- there is only one element in the stack, push the element back and push :error:

- the stack is empty, push :error: onto the stack

- if either of the top two elements are not integers, push back the elements and push :error:

For example:

```
┌─────────┐       ┌────────┐       ┌────────┐       ┌────────┐
│  input  │       │ stack  │       │ stack  │       │ stack  │
│ push 7  │   →   │   7    │   →   │   7    │   →   │ :true: │
│ push 7  │       └────────┘       │   7    │       └────────┘
│ equal   │                        └────────┘
└─────────┘
```

Consider another example:

```
┌──────────┐       ┌────────┐       ┌─────────┐       ┌─────────┐
│  input   │       │ stack  │       │  stack  │       │  stack  │
│ push 8   │   →   │   8    │   →   │ :error: │   →   │ :error: │
│ push 9.5 │       └────────┘       │   8     │       │ :error: │
│ equal    │                        └─────────┘       │   8     │
└──────────┘                                           └─────────┘
```

## 5.6   lessThan

The command lessThan refers to numeric less than ordering. This operator consumes the top two values on the stack and pushes the result (a single Boolean value) onto the stack. The :error: value will be pushed onto the stack if:

- there is only one element in the stack, push the element back and push :error:

- the stack is empty, push :error: onto the stack

- if either of the top two elements aren't integers, push back the elements and push :error:

For example:

| input |
|---|
| push 7 |
| push 8 |
| lessThan |

$\rightarrow$

| stack |
|---|
| 7 |

$\rightarrow$

| stack |
|---|
| 8 |
| 7 |

$\rightarrow$

| stack |
|---|
| :true: |

## 5.7   bind

The bind command binds a name to a value. It is evaluated by popping two values from the stack. The second value popped must be a name (see section on push for details on what constitutes a 'name'). The name is bound to the value (the first thing popped off the stack). The value can be any of the following:

- An integer

- A string

- Boolean

- :unit:

- The value of a name that has been previously bound

The name value binding is stored in an environment data structure. The result of a bind operation is :unit: which is pushed onto the stack. :error: will be pushed onto the stack if:

- If we are trying to bind an identifier to an unbound identifier, in which case all elements popped must be pushed back before pushing :error: onto the stack.

- the stack is empty, push :error: onto the stack.

### 5.7.1   Example 1

| input |
|---|
| push a |
| push 3 |
| bind |

$\rightarrow$

| stack |
|---|
| a |

$\rightarrow$

| stack |
|---|
| 3 |
| a |

$\rightarrow$

| stack |
|---|
| :unit: |

### 5.7.2  Example 2

```
input
push sum1
push 7          stack        stack        stack        stack        stack        stack
bind       →    sum1    →    7       →    :unit:  →    sum2    →    5       →    :unit:
push sum2                    sum1         :unit:       :unit:       sum2         :unit:
push 5                                                              :unit:
bind
```

You can use bindings to hold values which could be later retrieved and used by functionalities you already implemented. For instance in the example below, an addition on a + name1 in example1, would add 13 + 3 and push the result 16 onto the stack.

### 5.7.3  Example 3

```
input
push a
push 13
bind            stack        stack        stack        stack        stack        stack        stack
push name1 →    a       →    13      →    :unit:  →    name1   →    3       →    :unit:  →    a
push 3                       a            :unit:       :unit:       name1        :unit:       :unit:
bind                                                                :unit:                    :unit:
push a
push name1
add

        stack        stack
        name1        16
   →    a       →    :unit:
        :unit:       :unit:
        :unit:
```
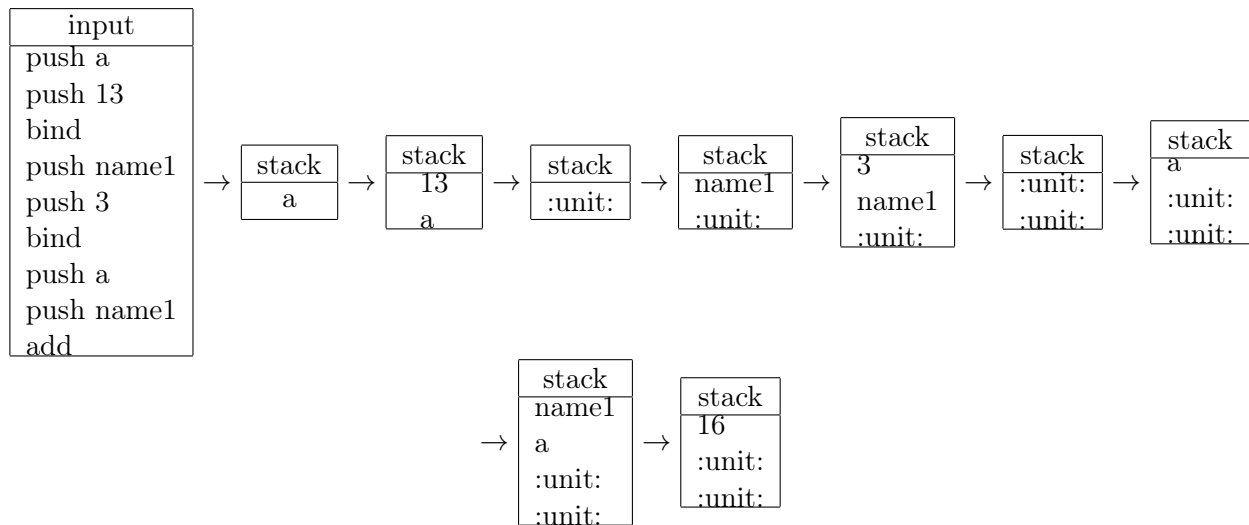
While performing operations, if a name has no binding, push :error: onto the stack, in which case all elements popped must be pushed back before pushing :error: onto the stack. Bindings can be overwritten, for instance:

```
input
push a
push 9
bind
push a
push 10
bind
```

Here, the second bind updates the value of 'a' to 10.

### Common Questions

(a)  What values can _name_ be bound to?

_name_ can be bound to integers, Boolean, string, :unit: and also previously bound values. For example,

1)

| input |
| --- |
| push a |
| push :true: |
| bind |

would bind a to :true:

2)

| input |
| --- |
| push a |
| 7.5 |
| bind |

would result in bind producing an :error: because a CANNOT be bound to :error:

3)

| input |
| --- |
| push b |
| let |
| push a |
| push 7 |
| bind |
| end |

would bind a to 7 and b to :unit:

4)

| input |
| --- |
| push b |
| push 8 |
| bind |
| push a |
| push b |
| bind |

would bind b to 8 and would bind a to the VALUE OF b which is 8.

5)

| input |
| --- |
| push b |
| push a |
| bind |

would result in an :error: because you are trying to bind b to an unbound variable a.

(b) How can we bind identifiers to previously bound values?

```
input
push a
push 7
bind
push b
push a
bind
```

The first bind binds the value of a to 7. The second bind statement would result in the name b getting bound to the VALUE of a—which is 7. This is how we can bind identifiers to previously bound values. Note that we are not binding b to a—we are binding it to the VALUE of a.

(c) Can we have something like this:

```
input
push a
push 15
push a
```

Yes. In this case 'a' is not bound to any value yet. And the stack contains:

```
stack
a
15
a
```

If we had :

```
input
push a
push 15
bind
push a
```

The stack would be :

```
stack
a
:unit:
```

(d) Can we push the same _name_ twice to the stack? For instance , what would be the result of the following:

```
input
push a
push a
quit
```

This would result in the following stack output:

14

```
stack
  a
  a
```

Yes, you can push the same _name_ twice to the stack. Consider binding it this way :

```
input
push a
push a
push 2
bind
```

This would result in
:unit: → as a result of binding a to 2
a → as a result of pushing the first a to the stack

(e) Output of the following code:

```
input
push a
push 9
bind
push a
push 10
bind
```

This would result in the following stack output:

would result in
:unit: → as a result of second bind
:unit: → as a result of first bind

## 5.8   if

The if command pops three values off the stack; x,y and z. The third value popped (z, in this case) must always be a Boolean. If z is :true:, executing the if command will push x back onto the stack, and if z is :false:, executing the if will push y back onto the stack.
:error: will be pushed onto the stack if:

- the third value is not Boolean, all elements (x,y, and z) should be pushed back onto the stack before pushing :error: onto the stack.

- the stack is empty, push :error: onto the stack

- there are less than 3 values on the stack in which case all elements popped must be pushed back before pushing :error: onto the stack.

For example:

```
input
push :true:        stack       stack      stack      stack
push 9       →     :true:   →   9      →   8      →   8
push 8                           :true:     9
push 8                                      :true:
if
```

**Common Questions**

(a) What values can 'if' take?

The result of executing a 'if' can be an integer or Boolean or string or :error: or :unit:

For instance,

|    | input         |
|----|---------------|
|    | push :true:   |
| a) | push "oracle" |
|    | push "jive"   |
|    | if            |

the result of if would be jive

|    | input        |
|----|--------------|
|    | push :false: |
|    | let          |
|    | push a       |
| b) | push 8       |
|    | bind         |
|    | end          |
|    | push 8.9     |
|    | if           |

the result of if would be :unit:

(b) What is the result of executing the following:

| input        |
|--------------|
| push a       |
| push 5       |
| bind         |
| pop          |
| push :true:  |
| push 4       |
| push a       |
| if           |

The stack would have a. Although the value of a is bound to 5, we only resolve the name to the value if we need to perform computation. (For 'if', the only value needed for computation is Boolean.)

## 5.9   let...end

let...end limits the scope of variables. "let" marks the beginning of a new environment—which is basically a sequence of bindings. The result of the let...end is the last stack frame of the let. Let...end can contain any number of operations but it will always result in a stack frame that is strictly larger than the stack prior to the let.

Trying to access an element that is not in scope of the let...end block would push :error: on the stack. let...end blocks can also be nested.

For example,

| input |
| --- |
| let |
| push c |
| push 13 |
| bind |
| let |
| push a |
| push 3 |
| bind |
| push a |
| push c |
| add |
| end |
| let |
| push b |
| push "ron" |
| bind |
| end |
| end |

## Original Stack

### 1st Let Expression

| stack |
|-------|
| c |

$\rightarrow$

| stack |
|-------|
| 13 |
| c |

$\rightarrow$

| stack |
|-------|
| :unit: |

$\rightarrow$

#### 2nd Let Expression

| stack |
|-------|
| a |
| :unit: |

$\rightarrow$

| stack |
|-------|
| 3 |
| a |
| :unit: |

$\rightarrow$

| stack |
|-------|
| :unit: |
| :unit: |

$\rightarrow$

| stack |
|-------|
| a |
| :unit: |
| :unit: |

$\rightarrow$

| stack |
|-------|
| c |
| a |
| :unit: |
| :unit: |

$\rightarrow$

| stack |
|-------|
| 16 |
| :unit: |
| :unit: |

$\rightarrow$

| stack |
|-------|
| 16 |
| :unit: |

$\rightarrow$

#### 3rd Let Expression

| stack |
|-------|
| b |
| 16 |
| :unit: |

$\rightarrow$

| stack |
|-------|
| ron |
| b |
| 16 |
| :unit: |

$\rightarrow$

| stack |
|-------|
| :unit: |
| 16 |
| :unit: |

$\rightarrow$
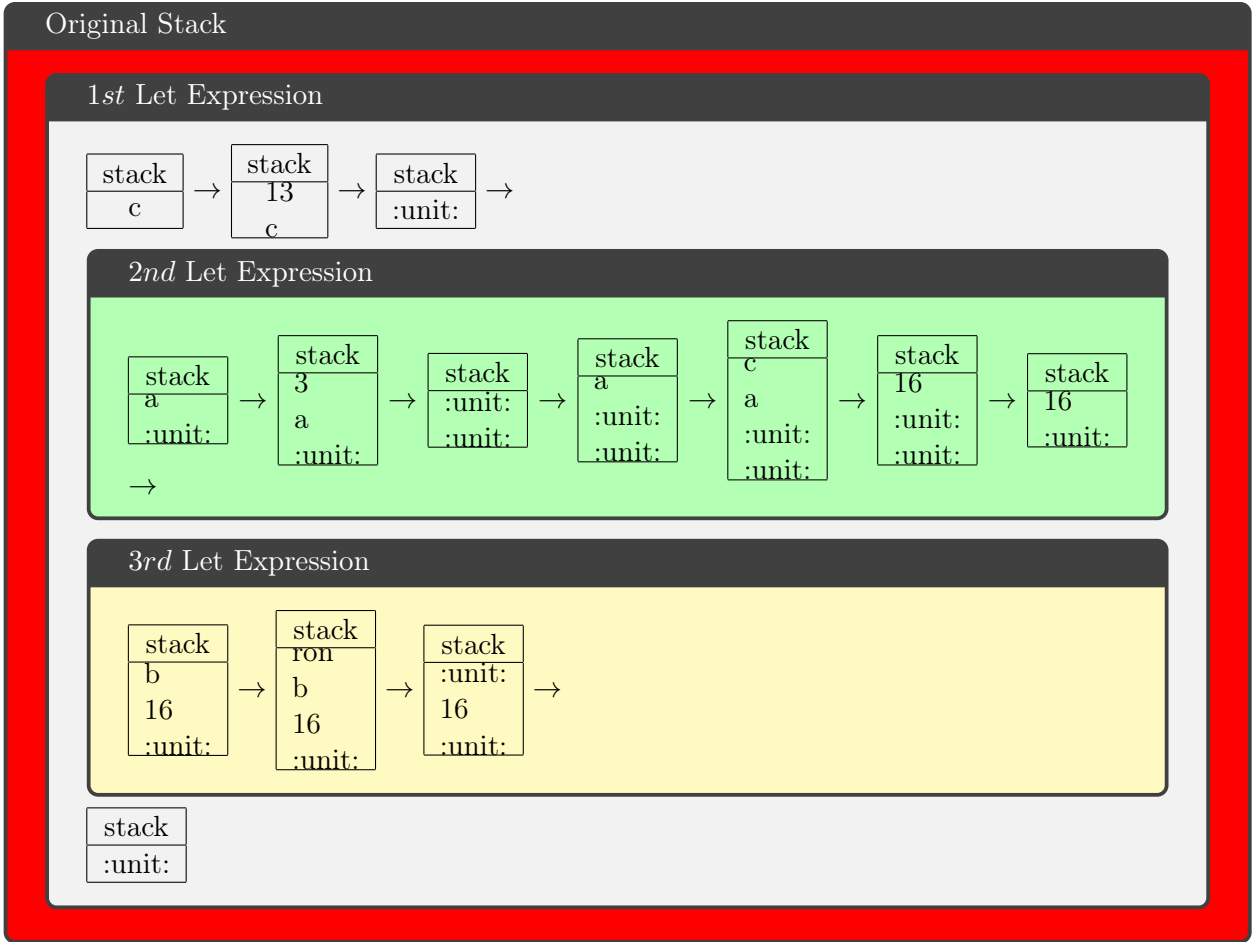
| stack |
|-------|
| :unit: |

## Common Questions

(a) What would be the output of running the following :

| input |
|-------|
| push 1 |
| let |
| push 2 |
| push 3 |
| push 4 |
| end |
| push 5 |

This would result in the stack :

| stack |
|-------|
| 5 |
| 4 |
| 1 |

Explanation : After the let...end is executed the last frame is returned—which is why we have 4 on the stack.

18

(b) What would be the result of executing the following :

| input |
|-------|
| let |
| push a1 |
| push 7.2 |
| bind |
| end |
| quit |

7.2 cant be pushed to the stack and a1 cannot be bound to :error: so, the result would be :error:

(c) What would be the output of running the following code:

| input |
|-------|
| let |
| push 3 |
| push 10 |
| end |
| add |
| quit |

The stack output would be

| stack |
|-------|
| :error: |
| 10 |

In the above example, the first let statement creates an empty environment (environment 1), then the name c is bound to 13. The result of this bind is a :unit: on the stack and a name value pair in the environment. The second let statement creates a second empty environment. Name a is bound here. To add a and c, these names are first looked up for their values in the current environment. If the value isn't found in the current environment, it is searched in the outer environment. Here, c is found from environment 1. The sum is pushed to the stack. A third environment is created with one binding 'b'.The second last end is to end the scope of environment 3 and the last end statement is to end the scope of environment 1. You can assume that the stack is left with at least 1 item after the execution of any let...end block.

# 6 Part 3: Functions  Due: Tuesday December 4 2018, at 11:59 pm

## 6.1 Functions

fun name1 name2

Denotes a function declaration, i.e. the start of a function called name1, which has one formal parameter name2. The expressions that follow comprise the function body. The function body is terminated with a special keyword funEnd. Note, name1 and name2 can be any valid name, but will never be any of the keywords in our language (e.g. add, push, pop, fun, funEnd, etc.). Also the function name and argument name cannot be the same.

<center>funEnd</center>

Denotes the end of a function body

<center>push funName</center>
<center>push arg</center>
<center>call</center>

Denotes applying the function funName to the actual parameter arg. When call is evaluated, it will apply the function funName to arg and pop both funName and arg from the stack. arg can either be a name (this includes function names), an integer, a string, boolean, or :unit:. :error: is pushed on the stack if either funName and arg are not bound in the current environment or if funName is not bound to a closure in the current environment. :error: is also pushed if the stack size is less than 2 when evaluating call.

When the interpreter encounters a function declaration expression it should being construction a closure. A closure will consist of (1) an environment, (2) the code for the function (the expressions between the function declaration and funEnd), and (3) the name of the formal parameter. :unit: should be pushed to the stack once the function declaration is evaluated and the closure created and bound to the function name in the environment.

1. The environment for the closure will be a copy of the current environment. (Challenge: if you would like to optimize your closure representation you do not need the entire environment, just the bindings of the variables used inside the function that are not defined inside the function and are not the formal parameter).

2. To compute the code for the function, you should copy all the expressions in order starting with the first expressions after the function declaration up to, but not including the funEnd.

3. In the current environment you should created a binding between the function name and its closure.

When a function is called, you should first check to see if there is a binding in the current environment, which maps funName to a closure. If one does not exists push :error: onto the stack. You should then check to see if the current environment contains a binding for arg, if it is a name instead of a value. If it does not then you should push :error: onto the stack. If arg is an :error: you should push :error: onto the stack.

If both funName and arg have appropriate bindings, or arg is a valid value, then the call to the function can proceed. To do this push the environment stored in the closure onto the stack. To this environment add a binding between the formal parameter (extracted from the closure) and the value of the actual parameter (arg). Note that if arg is a name, then it will have a binding in the environment at the point of the call (i.e. the environment before you pushed the environment stored in the closure). You should then save the current stack and create a new stack that will be used for the execution of the function (note: you may want to implement the stack as a stack of stacks to handled nested function calls and recursion, much like implementing the environment as a stack of maps). Next retrieve the code for the function and begin executing the expressions. The function completes once the last expression in code for the function is executed. When this happens you should restore the environment to the environment that existed prior to the function call (Hint: if you are implementing your environment as a stack of local environments, this will entail popping of the top environment.). The stack should also be restored to what the stack was at the point of the call (hint: if you implemented your stack as a stack of stacks, this only requires popping of the top

<center>20</center>

stack to restore the stack to what it was prior to the call). Once the environment has been restored, execution should resume with the expression that follows the call.

<div align="center">return</div>

Functions can return values by using a return expression. Since functions themselves are values (a closure), this means functions can take other functions as arguments and can return functions. When a return expression is evaluated, the function stops execution. When this happens you should restore the environment to the environment that existed prior to the function call, just like if the function completed by execution the last expression in the function's code. The stack should also be restored to what the stack was at the point of the call. Additionally you should push the last stack frame the function pushed onto the restored stack (the stack at the point of the call).
Please note that background color and indentation is used only to improve readability. Closure would consist of code within colored background.

### 6.1.1 Example 1

| input |
| --- |
| fun identity x |
| push x |
| return |
| funEnd |
| push identity |
| push 1 |
| call |
| quit |

→

| stack |
| --- |
| 1 |
| :unit: |

1 → return value of calling identity and passing in x as an argument
:unit: → result of declaring identity

### 6.1.2 Example 2

| input |
| --- |
| fun identity x |
| push x |
| return |
| funEnd |
| push identity |
| push 1.2 |
| call |
| quit |

→

| stack |
| --- |
| :error: |
| :error: |
| identity |
| :unit: |

:error: → error as a result of calling a function with error as the actual parameter
:error: → result of pushing 1.2
identity → push of identity
:unit: → result of declaring identity

<div align="center">21</div>

### 6.1.3    Example 3

| input |
| --- |
| <div style="background:yellow">fun identity x<br>push x<br>return<br>funEnd</div> |
| push x |
| push 1 |
| bind |
| push identity |
| push x |
| call |
| quit |

→

| stack |
| --- |
| 1 |
| :unit: |
| :unit: |

1 → return value of calling identity and passing in x as an argument
:unit: → result of binding x
:unit: → result of declaring identity

### 6.1.4    Example 4

| input |
| --- |
| push x |
| push 3 |
| bind |
| <div style="background:yellow">fun addX arg<br>push x<br>push arg<br>add<br>return<br>funEnd</div> |
| push x |
| push 5 |
| bind |
| push a |
| push 3 |
| bind |
| push addX |
| push a |
| call |
| quit |

→

| stack |
| --- |
| 6 |
| :unit: |
| :unit: |
| :unit: |
| :unit: |
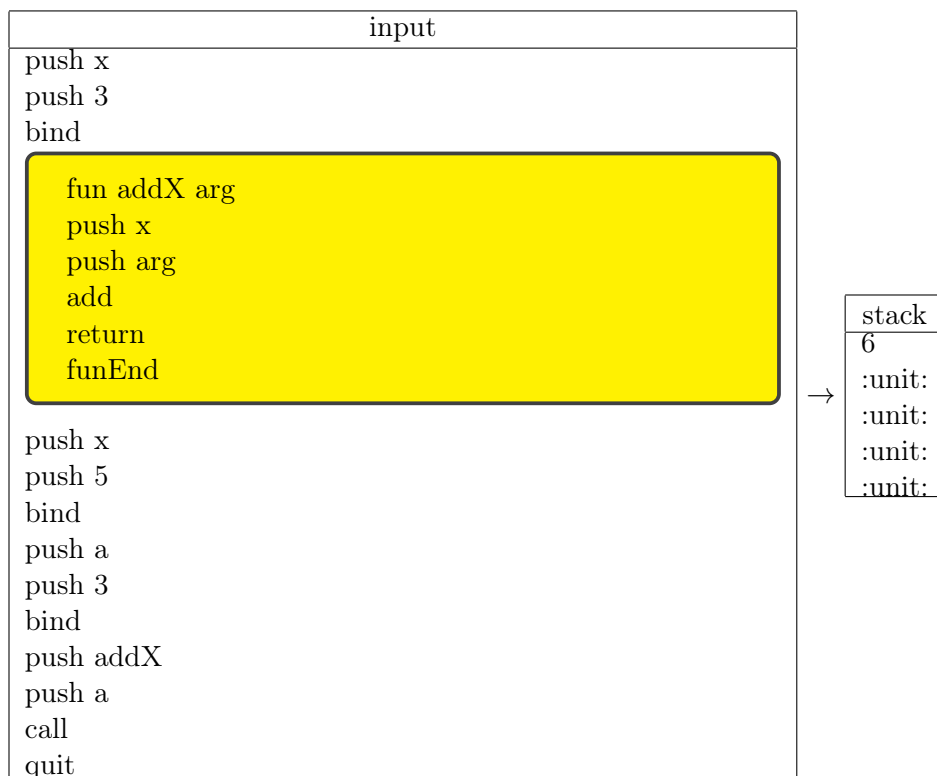
6 → result of function call
:unit: → result of third binding
:unit: → result of second binding
:unit: → result of function declaration
:unit: → result of first binding

### 6.1.5   Example 5

```
                          input
  ┌──────────────────────────────────────────┐
  │ fun stop arg                             │
  │ push 1                                   │
  │ return                                   │
  │ funEnd                                   │
  └──────────────────────────────────────────┘
  ┌──────────────────────────────────────────┐
  │ fun factorial arg                        │
  │ push arg                                 │
  │ push 1                                   │
  │ sub                                      │
  │ push 1                                   │
  │ push arg                                 │
  │ equal                                    │
  │ push factorial                           │
  │ push stop                                │
  │ if                                       │
  │ swap                                     │
  │ call                                     │
  │ push arg                                 │
  │ mul                                      │
  │ return                                   │
  │ funEnd                                   │
  └──────────────────────────────────────────┘
   push factorial
   push 3
   call
   quit
```

| stack |
|-------|
| 6     |
| :unit: |
| :unit: |

$\rightarrow$

6 $\rightarrow$ value returned from factorial

:unit: $\rightarrow$ declaration of factorial

:unit: $\rightarrow$ declaration of stop

### 6.1.6 Example 6

```
                          input
  ┌──────────────────────────────────────────┐
  │ fun add1 x                                 │
  │ push x                                     │
  │ push 1                                     │
  │ add                                        │
  │ return                                     │
  │ funEnd                                     │
  └──────────────────────────────────────────┘
  push z
  push 2
  bind
  ┌──────────────────────────────────────────┐        ┌─────────┐
  │ fun twiceZ y                               │        │  stack  │
  │ push y                                     │        ├─────────┤
  │ push z                                     │        │    6    │
  │ call                                       │    →   │ :unit:  │
  │ push y                                     │        │ :unit:  │
  │ push z                                     │        │ :unit:  │
  │ call                                       │        └─────────┘
  │ push y                                     │
  │ push z                                     │
  │ call                                       │
  │ add                                        │
  │ return                                     │
  │ funEnd                                     │
  └──────────────────────────────────────────┘
  push twiceZ
  push add1
  call
  quit
```
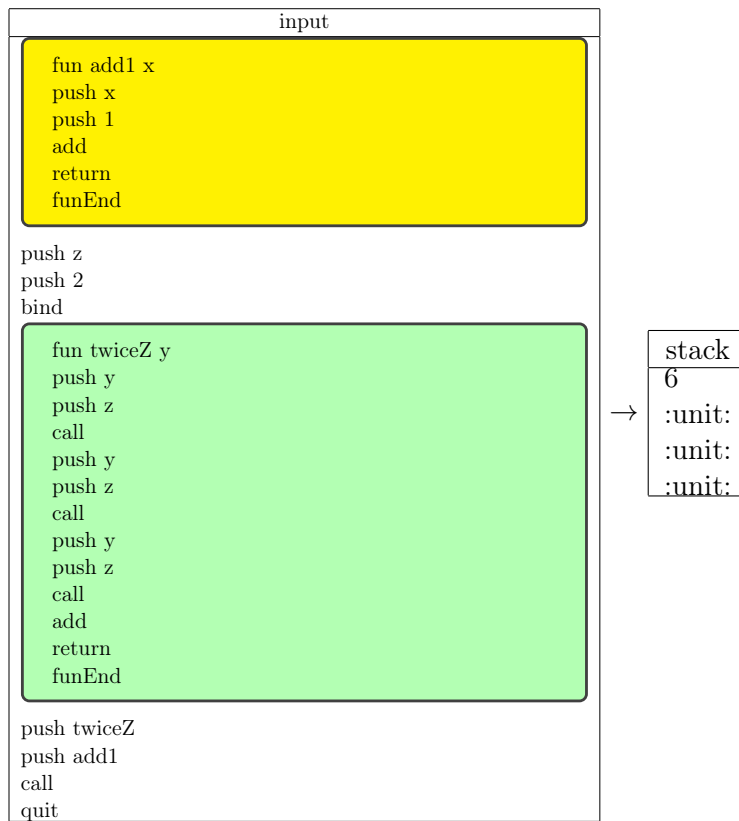
6 → return of calling twiceZ and passing add1 as an argument

:unit: → declaration of twiceZ

:unit: → binding of z

:unit: → declaration of the add1 function

## 6.2 Functions and Let

Functions can be declared inside a Let expression. Much like the lifetime of a variable binding, the binding of a function obeys the same rules. Since Let introduces a stack of environments, the closure should also take this into account. The easiest way to implement this is for the closure to store the stack of environments present at the declaration of the function. (note: you can create a more optimal implementation by only storing the bindings of the free variables you for the function—to do this you would look up each free variable in the current environment and add a binding from the free variable to the value in the environment stored in the closure)

(please note background color is used only to improve readability):

### 6.2.1  Example 1

```
input
┌──────────────────────────────────────────────┐
│  let                                          │
│   ┌────────────────────────────────────────┐  │
│   │   fun identity x                       │  │              stack
│   │   push x                               │  │           ┌──────────┐
│   │   return                               │  │           │  :error: │
│   │   funEnd                               │  │    →      │  1       │
│   └────────────────────────────────────────┘  │           │  identity│
│  end                                          │           │  :unit:  │
└──────────────────────────────────────────────┘           └──────────┘
push identity
push 1
call
quit
```

:error: → error since identity is not bound in the environment

1 → push of 1

identity → push of identity

:unit: → result of declaring identity, this is the result of the Let expression

### 6.2.2  Example 2

```
input
┌──────────────────────────────────────────────┐
│  fun identity x                               │
│   ┌────────────────────────────────────────┐  │
│   │   let                                  │  │              stack
│   │   push x                               │  │           ┌──────────┐
│   │   end                                  │  │           │  1       │
│   └────────────────────────────────────────┘  │    →      │  :unit:  │
│  return                                       │           └──────────┘
│  funEnd                                       │
└──────────────────────────────────────────────┘
push identity
push 1
call
quit
```
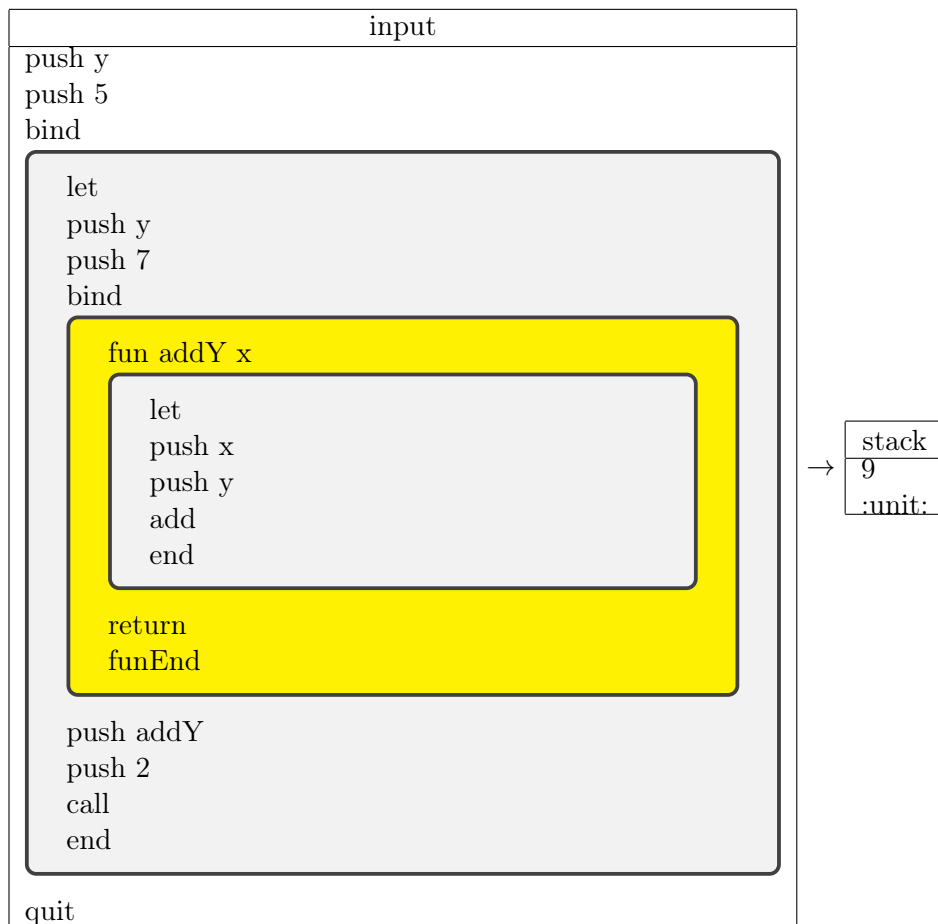
1 → return value of calling identity and passing in x as an argument

:unit: → result of declaring identity

### 6.2.3 Example 3

| input |
|---|
| fun double x |
|     let |
|     push x |
|     push x |
|     add |
|     end |
|   return |
|   funEnd |
| push double |
| push 2 |
| call |
| quit |

$\rightarrow$

| stack |
|---|
| 4 |
| :unit: |

4 $\rightarrow$ return value of calling identity and passing in x as an argument

:unit: $\rightarrow$ result of declaring identity

### 6.2.4 Example 4

| input |
|---|
| push y |
| push 5 |
| bind |
|   let |
|   push y |
|   push 7 |
|   bind |
|     fun addY x |
|       let |
|       push x |
|       push y |
|       add |
|       end |
|     return |
|     funEnd |
|   push addY |
|   push 2 |
|   call |
|   end |
| quit |

$\rightarrow$
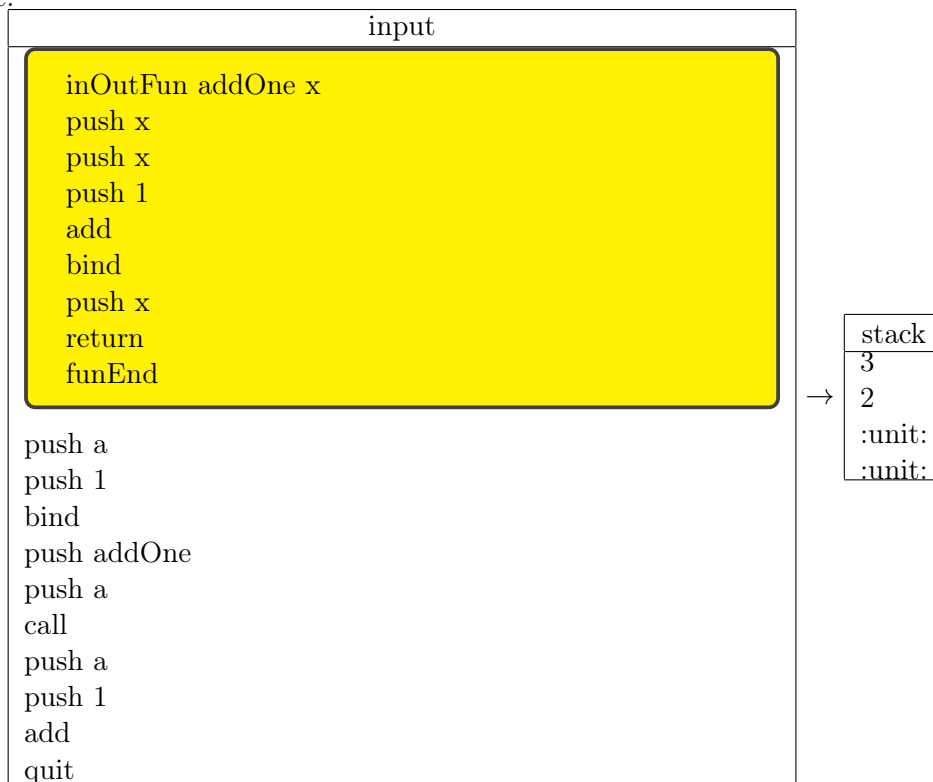
| stack |
|---|
| 9 |
| :unit: |

26

9 → return value of calling identity and passing in 2 as an argument
:unit: → result of binding y to 5

## 6.3   In/Out Functions

Our language will also support in/out parameters for specially denoted functions. Instead of using the fun keyword, functions that have in/out parameters are declared using the inOutFun keyword. In/out functions behave just like regular functions and all the rules defined for functions apply. In addition, when an in/out function returns, the value bound to the formal parameter is bound to the actual parameter in the environment after the call.

In/out functions should have a similar implementation to regular functions. To this implementation you should add an additional operation when the function returns. In addition to restoring the environment at the call site, the return will do a look up of formal parameter in the environment for the function. This value will be bound to the actual parameter in the environment at the call site.

| input |
|---|
| inOutFun addOne x |
| push x |
| push x |
| push 1 |
| add |
| bind |
| push x |
| return |
| funEnd |
| push a |
| push 1 |
| bind |
| push addOne |
| push a |
| call |
| push a |
| push 1 |
| add |
| quit |

→

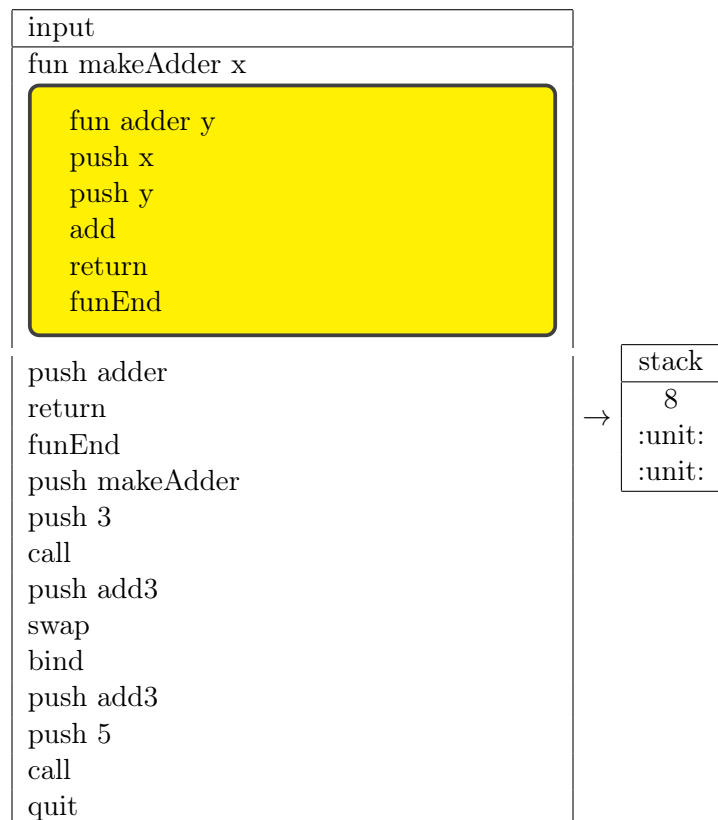| stack |
|---|
| 3 |
| 2 |
| :unit: |
| :unit: |

3 → result of add (note a is bound to two)
2 → return value of calling addOne and passing in x as an argument
:unit: → result of binding a
:unit: → result of declaring addOne

## 6.4 First-Class Functions

This language treats functions as any other value. They can be used as arguments to functions, and can be returned from functions.
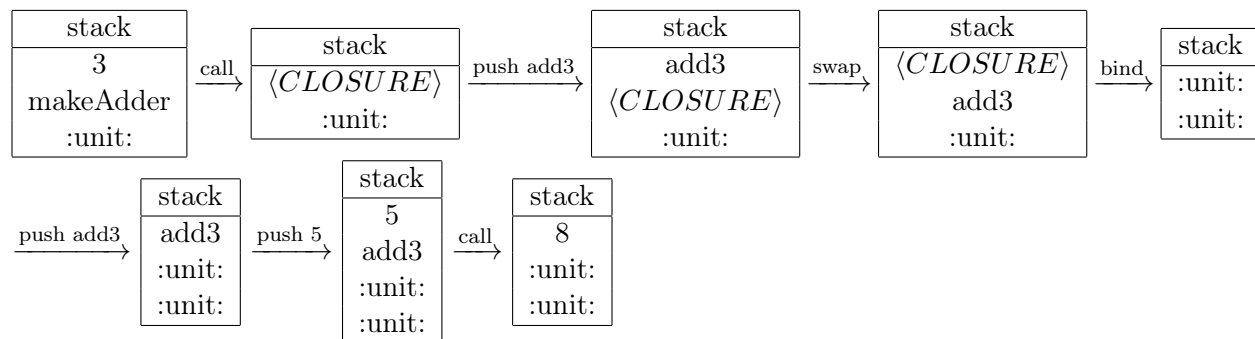
### 6.4.1 Example 1: Curried adder

| input |
| --- |
| fun makeAdder x |

```
        fun adder y
        push x
        push y
        add
        return
        funEnd

    push adder
    return
    funEnd
    push makeAdder
    push 3
    call
    push add3
    swap
    bind
    push add3
    push 5
    call
    quit
```

$\rightarrow$

| stack |
| --- |
| 8 |
| :unit: |
| :unit: |

$8 \rightarrow$ Evaluated from calling the generated function add3 with argument 5
:unit: $\rightarrow$ The result of binding the generated function to the name add3
:unit: $\rightarrow$ The result of declaring the function makeAdder

Step by step (after declaring makeAdder, pushing 3, and pushing makeAdder):

| stack |
| --- |
| 3 |
| makeAdder |
| :unit: |

$\xrightarrow{\text{call}}$

| stack |
| --- |
| $\langle CLOSURE \rangle$ |
| :unit: |

$\xrightarrow{\text{push add3}}$

| stack |
| --- |
| add3 |
| $\langle CLOSURE \rangle$ |
| :unit: |

$\xrightarrow{\text{swap}}$

| stack |
| --- |
| $\langle CLOSURE \rangle$ |
| add3 |
| :unit: |

$\xrightarrow{\text{bind}}$

| stack |
| --- |
| :unit: |
| :unit: |

$\xrightarrow{\text{push add3}}$

| stack |
| --- |
| add3 |
| :unit: |
| :unit: |

$\xrightarrow{\text{push 5}}$

| stack |
| --- |
| 5 |
| add3 |
| :unit: |
| :unit: |

$\xrightarrow{\text{call}}$

| stack |
| --- |
| 8 |
| :unit: |
| :unit: |

If a function is returned from another function, it need not be bound to a name in the environment it is returned in. For example:

| input |
|---|
| fun identity x |
| push x |
| return |
| funEnd |
| fun _catExcl y |
| push y |
| push "!" |
| cat |
| return |
| funEnd |
| push identity |
| push _catExcl |
| call |
| push "Dunder Mifflin" |
| call |
| quit |

$\rightarrow$

| stack |
|---|
| Dunder Mifflin! |
| :unit: |
| :unit: |

Dunder Mifflin! $\rightarrow$ Computed from calling the *closure* returned by the identity function applied to concatExcl with the argument "Dunder Mifflin".

:unit: $\rightarrow$ The result of declaring the function _catExcl.

:unit: $\rightarrow$ The result of declaring the identity function.

Here is a closer look at how the stack develops through this program. Note that function closures will never be on the stack when the program finishes execution.

| stack |
|---|
| concatExcl |
| identity |
| :unit: |
| :unit: |

$\xrightarrow{\text{call}}$

| stack |
|---|
| $\langle CLOSURE \rangle$ |
| :unit: |
| :unit: |

$\xrightarrow{\text{push "Dunder Mifflin"}}$

| stack |
|---|
| Dunder Mifflin! |
| $\langle CLOSURE \rangle$ |
| :unit: |
| :unit: |

$\xrightarrow{\text{call}}$

| stack |
|---|
| Dunder Mifflin! |
| :unit: |

1. You can make the following assumptions:

   - Expressions given in the input file are in correct formats. For example, there will not be expressions like "push", "3" or "add 5" .

   - No multiple operators in the same line in the input file. For example, there will not be "pop pop swap", instead it will be given as

     $$\text{pop}$$
     $$\text{pop}$$
     $$\text{swap}$$

   - No function closures will be left on the stack.

   - All `let` commands will have a matching `end`.

   - There will always be at least one value inside the final stack.

   - There will always be a "quit" in the input file to exit your interpreter and output the stack.

2. You can assume that all test cases will have a quit statement at the end.

3. You can assume that your interpreter function will only be called ONCE per execution of your program.

# Step by step examples

1. If your interpreter reads in expressions from "input", states of the stack after each operation are shown below:

| input |
|---|
| push 10 |
| push 15 |
| push 30 |
| sub |
| push :true: |
| swap |
| add |
| pop |
| neg |
| quit |

First, push 10 onto the stack:

| stack |
|---|
| 10 |

Similarly, push 15 and 30 onto the stack:

| stack |
|---|
| 30 |
| 15 |
| 10 |

sub will pop the top two values from the stack, calculate 15-30 = -15, and push -15 back:

| stack |
|---|
| -15 |
| 10 |

Then push the boolean literal :true: onto the stack:

| stack |
|---|
| :true: |
| -15 |
| 10 |

swap consumes the top two values, interchanges them and pushes them back:

30

| stack |
|-------|
| -15 |
| :true: |
| 10 |

add will pop the top two values out, which are -15 and :true:, then calculate their sum. Here, :true: is not a numeric value therefore push both of them back in the same order as well as an error literal :error:

| stack |
|-------|
| :error: |
| -15 |
| :true: |
| 10 |

pop is to remove the top value from the stack, resulting in:

| stack |
|-------|
| -15 |
| :true: |
| 10 |

Then after calculating the negation of -15, which is 15, and pushing it back, quit will terminate the interpreter and write the following values in the stack to *outputFile*:

| stack |
|-------|
| 15 |
| :true: |
| 10 |

Now, go back to the example inputs and outputs given before and make sure you understand how to get those results.

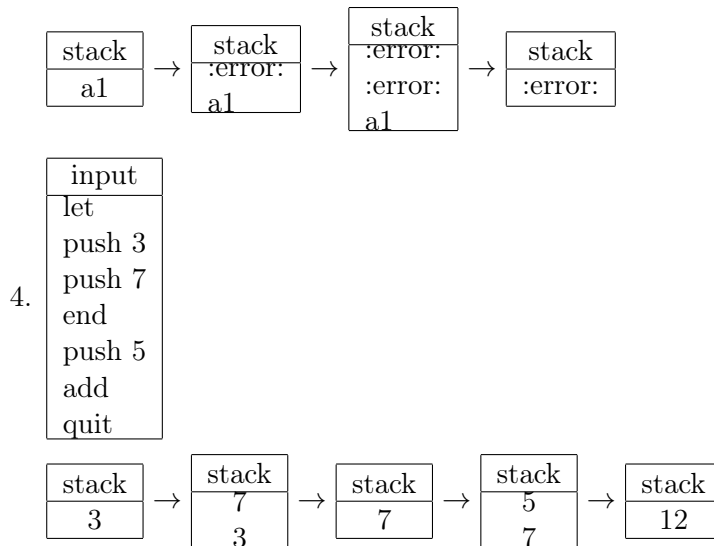2. More Examples of bind and let...end:

| input |
|-------|
| push a |
| push 17 |
| add |

| stack |
|-------|
| a |

$\rightarrow$

| stack |
|-------|
| 17 |
| a |

$\rightarrow$

| stack |
|-------|
| :error: |
| 17 |
| a |

The error is because we are trying to perform an addition on an unbound variable "a".

3.

| input |
|-------|
| let |
| push a1 |
| push 7.2 |
| bind |
| end |

| stack |
|-------|
| a1 |

$\rightarrow$

| stack |
|-------|
| :error: |
| a1 |

$\rightarrow$

| stack |
|-------|
| :error: |
| :error: |
| a1 |

$\rightarrow$

| stack |
|-------|
| :error: |

4.

| input |
|-------|
| let |
| push 3 |
| push 7 |
| end |
| push 5 |
| add |
| quit |

| stack |
|-------|
| 3 |

$\rightarrow$

| stack |
|-------|
| 7 |
| 3 |

$\rightarrow$

| stack |
|-------|
| 7 |

$\rightarrow$

| stack |
|-------|
| 5 |
| 7 |

$\rightarrow$

| stack |
|-------|
| 12 |

Explanation :

Push 3

Push 7

Pushes 3 and 7 on top of the stack. When you encounter the "end", the last stack frame is saved (which is why the value of 7 is retained on the stack) , then 5 is pushed onto the stack and the values are added.

# 7 Frequently Asked Questions

1. Q: What are the contents of test case $X$?

   A: We purposefully withhold some test cases to encourage you to write your own test cases and reason about your code. You cannot test *every* possible input into the program for correctness. We will provide high-level overviews of the test cases, but beyond that we expect you to figure out the functionalities that are not checked with the tests we provide. But you can (and should) run the examples shown in this document!

2. Q: Why does my program run locally but fail on Autolab?

   A: Check the following:

   - Ensure that your program matches the types and function headers defined in section 2 on page 1.
   - Make sure that any testing code is either removed or commented out. If your program calls interpreter with input "input.txt", you will likely throw an exception and get no points.
   - *Do not submit testing code.*
   - If you submit a zip or tar file, the main interpreter file must be at the top level of the archive. When extracted, there must be a file named `interpreter.{java,py,ml}` in the current directory.
   - `stdout` and `stderr` streams are not graded. Your program must write to the output file specified by *outputFile* for you to receive points.

- *Close your input and output files.*
- Java specific:
  - (a) Remove all package declarations from your program. If the interpreter class is inside a package, your program will not compile with the grading scripts.
  - (b) The class and method *interpreter* must be public. Additionally, the method must be `static` and `void`.
- Python specific:
  - (a) No external modules are installed (eg numpy); stick to the standard library.
  - (b) The arguments to `interpreter` will not be command line arguments. They will be passed in by the function's arguments.
- OCaml specific:
  - (a) Core and any other external libraries are not available.
  - (b) Autolab only supports 4.02, so any features added after are unsupported.
  - (c) The *Sys* module is unavailable.

3. Q: Why doesn't Autolab give useful feedback?

   A: Autolab is strictly a grading tool to tell you how many test cases you passed and your total score. Test and debug your program locally before submitting to Autolab.

4. Q: Are there any runtime complexity requirements?

   A: Although having a reasonable runtime and space complexity is important, the only official requirement is that your program runs the test suite in less than three minutes.

5. Q: Is my final score the highest score I received of all my submissions?

   A: No. Your final score is only your most recent submission.

6. Q: What can I do if an old submission received a better grade than my most recent submission?

   A: You can always download any of your previous submissions. If the deadline is approaching, we suggest resubmitting your highest-scoring submission before Autolab locks.